

Character Consistency Mastery: Building Persistent AI Characters Across Platforms

Module 3: Stable Diffusion & LoRA Training

Deep Control Through Custom Models

Table of Contents

1. Introduction to Custom Model Training
2. Understanding LoRA Technology
3. Setting Up Your Training Environment
4. Dataset Preparation: The Foundation of Success
5. Training Your First Character LoRA
6. Advanced Training Techniques
7. Using LoRAs in Production
8. Combining Multiple LoRAs and Techniques
9. Troubleshooting Training Issues
10. Building a Professional Character Pipeline

Chapter 1: Introduction to Custom Model Training

Why Custom Models Change Everything

While reference-based methods offer good consistency, custom model training provides absolute control.

- **Reference-Based (e.g., Midjourney):** Achieves 70-85% consistency, depends on prompt precision, and requires a reference image for every generation.
- **Custom LoRA (Stable Diffusion):** Achieves 95-99% consistency, works with simple trigger words, is encoded in model weights, and is portable across platforms.

Custom training becomes essential for commercial productions, long-form content with recurring characters, and projects requiring extreme pose variations or style consistency. The time investment (initially a few hours) yields the return of infinite consistent generations and complete asset ownership.

Chapter 2: Understanding LoRA Technology

What LoRA Actually Is

LoRA (Low-Rank Adaptation) is a training technique that creates a small "patch" file (10-150MB) for a large AI model (5GB+). This patch efficiently teaches the model new concepts, like a specific character, style, or object, without altering the base model itself. These LoRA files can be mixed, matched, and weighted to achieve precise results.

LoRA Compatibility

- **Base Model:** LoRAs trained on a specific base model architecture (e.g., SD 1.5 or SDXL) will only work with models from that same architecture.
- **Cross-Platform Usage:** The same LoRA file can be used across various Stable Diffusion interfaces, including ComfyUI, Automatic1111, Forge, and InvokeAI, as well as in custom pipelines.

Chapter 3: Setting Up Your Training Environment

Hardware Requirements

- **Minimum Viable Setup:** GPU with 8GB+ VRAM (e.g., RTX 30-series), 16GB System RAM, 100GB storage.
- **Recommended Setup:** GPU with 12GB+ VRAM (e.g., RTX 40-series), 32GB System RAM, 500GB SSD.
- **Professional Setup:** GPU with 24GB+ VRAM (e.g., RTX 5090 or professional equivalent), 64GB System RAM, 2TB NVMe SSD.

Software Stack

- **Option 1: Kohya_ss GUI (Recommended for Beginners):** A user-friendly interface with preset configurations and integrated tools, though primarily Windows-focused.
- **Option 2: Automatic1111 with Training Extension:** Integrates training directly into a familiar generation UI, allowing for immediate testing.
- **Option 3: Command Line Training (Advanced):** Offers maximum control and performance but has a steep learning curve and requires manual configuration.

Environment Configuration

A crucial step is setting up the correct software environment, including the CUDA toolkit and PyTorch. It's essential to install the version of PyTorch that matches your specific, installed CUDA version.

```
Bash
# Check CUDA availability
nvidia-smi
```

Install appropriate PyTorch (example for a CUDA version)

```
pip install torch torchvision --index-url https://download.pytorch.org/whl/cu121
```

For systems with limited VRAM, memory can be optimized by enabling gradient checkpointing, using an 8-bit Adam optimizer, and reducing the training batch size to 1.

Chapter 4: Dataset Preparation - The Foundation of Success

The 90/10 rule applies here: **90% of LoRA quality comes from the dataset, and 10% comes from training settings.** Never rush this phase.

Image Requirements

- **Quantity:** 20-30 images is recommended for a good balance, with diminishing returns after 50+.
- **Resolution:** Minimum 1024x1024 for SDXL models. Varied aspect ratios are better than all square.
- **Quality:** Images must have a sharp focus on the subject, good lighting, and minimal compression artifacts.

Dataset Curation Strategies

Ensure variety across angles, expressions, and lighting. Actively avoid including other people, heavily filtered images, extreme crops, or low-resolution sources in your dataset.

Image Preprocessing & Captioning

1. **Preprocessing:** Use tools to remove distracting backgrounds, consistently crop and align your subject, and standardize resolution.
2. **Captioning:** This step is critical. Captions teach the AI what to associate with your character. A good strategy is to use a unique **trigger word** (e.g., `ohwx person`) followed by a simple, descriptive caption of the image content. Automated tools like BLIP can generate base captions, which you can then refine.

Chapter 5: Training Your First Character LoRA

Essential Training Parameters

- **Learning Rate:** Controls how quickly the model learns. A standard starting point is `1e-4` or `2e-4`.
- **Network Architecture (Rank/Dim):** Determines the LoRA's capacity for detail. A rank of `128` is a balanced choice.
- **Training Duration:** Calculated in steps. A common configuration is `20 images × 10 repeats × 10 epochs = 2000 steps`.

Step-by-Step Training Walkthrough

1. **Configure Training Script:** In your chosen tool (e.g., Kohya_ss), set your source model, dataset path, and training parameters.
2. **Initialize Training:** Run a pre-flight check to ensure GPU memory is clear and paths are correct, then launch the training process.
3. **Monitor Progress:** Watch the loss curve (it should decrease smoothly) and VRAM usage.
4. **Test Checkpoints:** Generate test images every 500-1000 steps to monitor for overfitting (where the AI perfectly recreates training images but loses flexibility).

Chapter 6: Advanced Training Techniques

- **Multi-Resolution Training:** Training on images of various resolutions makes the final LoRA more flexible.
- **Regularization:** Using a folder of generic "class images" (e.g., photos of a "person" or "woman") prevents the model from "forgetting" how to draw generic people, which improves the LoRA's compatibility.
- **Dropout:** Intentionally ignoring some captions or network connections during training prevents overfitting and improves generalization.
- **Token Weighting:** Emphasize important features in your captions, like `(green eyes:1.1)`.

Chapter 7: Using LoRAs in Production

Integration with Generation UIs

In interfaces like Automatic1111 or ComfyUI, you can load a LoRA and control its strength. The prompt syntax often looks like this:

```
a photo of <lora:character_v1:0.8> in a park
```

The number `(0.8)` is the weight, controlling the LoRA's influence. A weight of 0.7-0.9 is often the sweet spot.

Prompt Engineering with LoRAs

Use a unique trigger word to call your character in the prompt. You can also use negative prompts to prevent common LoRA artifacts.

Negative prompt: "multiple faces, distorted features, inconsistent appearance, style mixing"

Chapter 8: Combining Multiple LoRAs and Techniques

You can stack LoRAs to combine effects, such as a character, a style, and an outfit.

```
<lora:character:0.8> <lora:art_style:0.6> <lora:uniform:0.7>
```

For maximum consistency, combine your LoRA with other tools like **ControlNet** for pose control and **IPAdapter** for face similarity. This "triple lock" system can yield 99%+ consistency.

Chapter 9: Troubleshooting Training Issues

- **Overfitting:** The LoRA only recreates training images. **Solution:** Reduce training steps, add more dataset variety, or lower the learning rate.
- **Underfitting:** The LoRA has a weak resemblance to the character. **Solution:** Increase training steps, improve dataset quality, or slightly raise the learning rate.
- **Style Bleeding:** The LoRA learns the background or clothing instead of the face. **Solution:** Use highly varied backgrounds and clothing in your dataset.

Chapter 10: Building a Professional Character Pipeline

Production Standards

A professional pipeline requires strict organization. Create a dedicated folder structure for each character project to manage references, development assets, training files, and final LoRAs.

Documentation Standards

Create a "Character Bible" for each character. This document should include the trigger word, key features, optimal LoRA weight, known limitations, and example prompts. This is essential for team collaboration and long-term project consistency.

Automation Workflows

For professional use, automate the pipeline with scripts. You can create scripts to automatically preprocess datasets, launch training with specific configurations, and run a suite of tests on the generated LoRA to verify its quality and consistency.

Final Project Challenge

A final challenge to solidify these skills involves creating a complete character system from scratch. This includes designing a unique character, preparing a professional dataset of 30+ images, training an optimized LoRA, generating a diverse set of 50 images to test its capabilities, and documenting the entire process in a character bible.