

What This Does

You're building commercial-grade product shots that look like cinema. Motorcycle ads for example. Watch campaigns. Anything that needs film-quality stills but starts from zero.

The method: chain three models. **Nano Banana** blocks composition. **Seedream 4.0** renders 4K film finish. **Midjourney** adds photorealism. Then fix product accuracy in Photoshop.

Built for creatives who don't want to fight prompts or settle for "close enough."

Quick Start

- Pick a theme. Write 5–7 shots. One line each.
- Hunt layout references online. Composition matters. Subject doesn't.
- Freepik → **Google Nano Banana** → feed layout ref + prompt: "Replace [subject] with [your product], keep angle and scale."
- Re-feed your best frame. Prompt angle variations: wide, low, tight, hero.
- Switch to **Seedream 4.0** → upload Nano Banana frame + film still as grade reference → render 4K.
- Upload 4K stills to Midjourney as image prompts. Feed ChatGPT each frame → auto-generate 1000–2000 char text prompt (product/brand/color explicit). Combine: image + text + `--raw` + style ref.
- Iterate with Vary (Strong/Subtle). Upscale winners.
- Photoshop → select product area → Nano Banana plugin → fix brand details while keeping MJ realism everywhere else.

Done.

Why This Works

Most workflows rely on one model. One model means compromise: layout control *or* realism *or* brand accuracy. Never all three.

This workflow splits the work. Each model does one thing perfectly. Nano Banana follows instructions. Seedream gives you film finish. Midjourney brings photographic texture. Photoshop corrects product fidelity.

You control layout, grade, material response, and brand accuracy at each step. No guessing. No hoping.

The Workflow

1. Theme + Storyboard

Start with intent. One sentence. "Night-rain motorcycle spot." "Luxury watch, desert dawn." Whatever.

Write 5–7 shots. One line each. Every shot needs a job:

- Open
- Reveal
- Pursuit
- Intimacy
- Product hero
- Payoff

Example:

- Wide → bike alone, wet street, sodium lights, leading lines
- Low tracking → front wheel, ground level, spray and motion blur
- Tight macro → brake disc, glowing, shallow depth
- Hero 3/4 → front, eye level, headlight flare, rain streaks
- Detail → logo, wet chrome, reflections

Keep it visual. No story beats. Just frames.

2. Hunt Layout References

Google Images. Pinterest. Instagram. Doesn't matter.

For each shot, grab 3–4 frames that match the **layout vibe**:

- Camera height
- Negative space
- Motion lines
- Where the eye lands

Ignore color. Ignore subject. A car ad's framing works for a bike. A watch campaign's macro works for a sneaker. You're stealing composition, not copying the image.

Drop links and notes under each storyboard cell so you see rhythm and coverage before you render anything.

3. Block Shots with Nano Banana

Open **Freepik AI Image Generator**. Model dropdown: **Google Nano Banana (Gemini 2.5 Flash Image)**.^[1]

Add your layout reference under **Image Reference / Visual Prompts**.

Prompt:

"Replace the car with [motorcycle make/model, finish]. Keep camera angle, horizon height, and subject scale. Preserve reflections and rain streaks. Add micro-spray from rear tire, faint brake disc blur. Night, wet road, practical sodium lights, shallow depth, cinematic motion trails. No text."

Generate.

Why Nano Banana? It follows instructions. You say "replace X with Y, keep the rest" and it respects the frame structure. Fast. Versatile. Weak on material realism, but that's not its job.

Pick your strongest frame.

4. Expand Coverage

Re-feed that frame as the new reference. Remove the original layout.

Prompt angle variations:

- **Wide establishing** → "same scene, 24mm equivalent, full bike in frame, context visible"
- **Low tracking** → "same scene, 35mm, camera at ground level, dramatic perspective"
- **Tight macro** → "same scene, 85mm, focus on front wheel and brake disc, shallow depth"
- **Hero 3/4** → "same scene, 50mm, eye level, front 3/4 angle, headlight prominent"

You're building shot diversity on one scene while keeping object identity stable. If consistency drifts, lock a seed or re-feed the strongest variant.

Repeat for all storyboard cells.

5. 4K Film Finish with Seedream

Take every Nano Banana frame you like. Now make it look like cinema.

Switch model to **Seedream 4.0 (4K)** in Freepik (Generator or Image Editor).

Upload your Nano Banana frame as **base image**.

Optional but powerful: add a **second image** as grade reference. A film still with the palette you want. Blade Runner 2049 orange. Drive neon pink. Mad Max dust and steel. Whatever.

Prompt:

"Base: [Nano Banana frame]. Grade reference: [film still]. Recreate at 4K. Match color palette, black density, midtone contrast, highlight halation. Preserve composition, position, and perspective. Keep realistic material response on metal and rubber. No added props. No text."

Seedream renders up to 4K. Sharpens materials without looking plastic. Respects your composition while applying the grade. Fast iteration.

Why not just upscale Nano Banana? Because upscaling washes out the image. Makes it too clean. Seedream recreates at 4K with film-quality contrast and highlight roll-off. Different process. Better result.

Export your 4K stills.

6. Midjourney: Photorealism Layer

This is where it gets real.

Upload your 4K Seedream frames to **Midjourney** as **image prompts**.^[3]

But Midjourney needs two things: image + text.

Here's the trick. Feed each Seedream frame to **ChatGPT**. Prompt:

"Create a visual description of this image, 1000–2000 characters. Always mention product, brand, and color. Describe lighting, materials, camera angle, atmosphere. Be explicit about color palette. Mention the product every time it's visible. Stay visual. No metaphors. No 'as if' phrasing."

ChatGPT spits out a prompt. Copy it.

Now in Midjourney, combine:

- Image URL** (your Seedream 4K still)
- Text prompt** (from ChatGPT)
- `--raw` (turns off MJ auto-stylization, gives you control)
- `--sref [style reference image URL]` (optional: lock grade or film look across shots)^[5]
- `--sw [0-1000]` (style weight — higher = stronger style influence)

Example structure:

```
https://your-seedream-image.jpg [1000-char ChatGPT prompt] --raw --sref https://file-still.jpg --sw 500
```

Why this works: Seedream gives you 4K composition and film grade. Midjourney adds photographic micro-structure and material believability. Raw Mode stops MJ from "Midjourney-fying" the result. Your refs drive the look.

Generate. Upscale your favorite.

7. Iterate Fast

Midjourney gives you options:

- **Vary (Strong)** → major variations, new angles or lighting shifts
- **Vary (Subtle)** → small tweaks, preserve the frame
- **Remix Mode** → change the prompt while varying (adjust lighting, shift camera height, swap background)
- **Vary Region** → select an area, prompt only that zone

Use Vary (Strong) first. Get 3–4 alternatives per scene. Pick winners. Then Vary (Subtle) to refine.

If you have **Profiles** set up in Midjourney (saved style presets), activate them here. Keeps your look consistent across shots without re-typing style refs every time.

8. Photoshop: Product Correction

Midjourney makes things look real. But it's not built for brand accuracy. Your logo might shift. Product proportions might drift.

Fix:

- Open your final MJ image in **Photoshop**
- Select only the product area (lasso, mask, whatever)
- Use the **Nano Banana plugin** (Gemini 2.5 Flash Image integration)
- Feed your product reference
- Prompt only for the selected zone: "Replace with [exact product], match lighting and perspective, preserve surrounding context"

Result: MJ realism everywhere. Brand-accurate product exactly where it matters. Minimal quality loss because you're only regenerating the selected area.

Optional: Automate with Node Systems

If you're running multiple campaigns or need repeatable pipelines, build it as a node graph.

Tools: **Flora Fauna** or **Weavy**.

Chain the first steps:

- Nano Banana (layout swap)
- Seedream 4.0 (4K film finish)
- ChatGPT prompt generation
- Renderer (WAN 2.2 or FLUX.1 [pro])

Why not Midjourney? No API. But WAN 2.2 with strong Seedream refs gets you 90% of the way. FLUX.1 [pro] works too, especially if you feed it solid 4K bases and use Kontext mode for style locking.

For ultra-fidelity, you can inject a **3D model** of your product into WAN 2.2. Perfect dimensions. Perfect brand match. But that's advanced. Different module.

Model Roles (Plain Language)

Nano Banana (Gemini 2.5 Flash Image)

Job: Follow instructions. Swap subjects while preserving layout.
Strength: Versatile. Respects composition. Fast.
Weakness: Low material realism. Not 4K-ready.
When: Blocking shots. Angle exploration.

Seedream 4.0

Job: Render 4K film finish. Apply cinematic grade.
Strength: Sharp materials. Film-quality contrast and highlight response. Preserves composition.
Weakness: Slightly less "photographic" than MJ in texture detail.
When: After Nano Banana. Before Midjourney.

Midjourney (Raw Mode)

Job: Add photographic micro-structure and material believability.
Strength: Realistic texture. Gloss, grain, weathering. Lighting feels real.
Weakness: Doesn't follow brand specs perfectly. Can drift from exact product shape.
When: After Seedream. Before final product correction.

Nano Banana in Photoshop

Job: Fix product details without touching the rest.
Strength: Surgical. Preserves MJ realism everywhere else.
Weakness: Requires manual selection.
When: Final step. Only if product drifted.

Prompt Skeletons

Nano Banana (layout swap)

Reference: [layout image]. Replace [original subject] with [your product: make, model, finish]. Keep camera angle and scale.

Seedream 4.0 (4K + grade)

Base: [Nano Banana frame]. Grade reference: [film still]. Recreate at 4K. Match color palette, black density, midtone contrast, highlight halation. Preserve composition, position, and perspective. Keep realistic material response on metal and rubber. No added props. No text.

Midjourney (image + text + raw)

[Seedream 4K image URL] [ChatGPT-generated 1000–2000 char description] --raw --sref [style reference URL] --sw [style weight]

Nano Banana in Photoshop (product fix)

Replace selected area with [exact product: brand, model, finish]. Match surrounding lighting and perspective.

Troubleshooting

Nano Banana ignores my layout reference

→ Be more explicit in the prompt. "Keep camera angle at 15° below horizon. Preserve subject scale at 60% of frame width."

→ Lock a seed if you're iterating on one frame.

Seedream changes my composition

→ Make sure you're uploading the Nano Banana frame as **base image**, not just as a reference.

→ Explicitly state "preserve composition, position, and perspective" in the prompt.

Midjourney output looks too "Midjourney"

→ Use `--raw`. Always.

→ Feed stronger image refs. Your Seedream 4K base should be the primary driver.

→ Lower `--sw` if style ref is overpowering the image prompt.

Product details drift in Midjourney

→ Expected. That's why you fix it in Photoshop with Nano Banana.

→ Or skip MJ and use WAN 2.2 / FLUX.1 in a node system with tighter i2i control.

Colors don't match across shots

→ Use the same grade reference (film still) for all Seedream passes.

→ In Midjourney, lock the same `--sref` across the sequence.

4K Seedream output is too clean

→ Add subtle grain in post. Or include "film grain, subtle halation" in the Seedream prompt.

→ Don't over-sharpen. Cinema isn't clinically sharp.

FAQ

Can I do this without Freepik?

Yes. Nano Banana is available elsewhere. Seedream 4.0 is ByteDance's model, accessible through other platforms. But Freepik bundles both in one interface, which is why it's faster.

Why not just use Midjourney from the start?

Because MJ doesn't let you control layout as tightly. You'd spend 50 prompts trying to get the right angle. Nano Banana with a layout ref gets you there in one shot.

Do I need Photoshop?

Only if product accuracy matters (logos, proportions, brand colors). If you're making concept art or mood boards, skip it.

What if I don't have a film still for grading?

Find one. Google "[film name] color palette stills." Or use a LUT reference. Or skip the grade ref and let Seedream render neutral — you can grade in post.

Can I use this for video?

The stills, yes. Feed your hero frames to **PlxVerse V5** (image-to-video) inside Freepik for 5–8 second motion previews. Or use Runway, Pika, whatever. But full video production is a different workflow.

How long does this take?

Per shot: 2–5 minutes in Nano Banana, 1–2 minutes in Seedream, 3–5 minutes in Midjourney, 2 minutes in Photoshop. Call it 10–15 minutes per final frame once you're fast.

What about automation?

Flora Fauna or Weavy. Chain Nano Banana → Seedream → WAN 2.2. Save as a blueprint. Run batches. But you lose Midjourney's texture quality unless you manually feed MJ at the end.

Why WAN 2.2 or FLUX.1 in automation?

No Midjourney API. WAN 2.2 is solid with strong refs. FLUX.1 [pro] with Kontext mode locks style well. Both integrate into node systems.

Can I inject a 3D model?

Yes, but that's advanced. Load a 3D render into WAN 2.2 as a depth or control map. Perfect dimensions. Perfect brand match. Different module.